

Reliability experiment: $5 \times 14 \times 200 = 14,000$ episodes per method

5 training seeds

42, 77, 123, 314, 2718

14 scenarios

held-out, stress, OOD

200 episodes

per seed and scenario

Seed-level inference

paired effects + t-CI

Acceptance gate

speed and reliability

Latency benchmark: batch 1 · warm-up 50 · 2,000 repeats · synchronized end-to-end timing