

GLOBE++ Lite-P+: Global-to-Local Policy Distillation with Predictive Risk Switching for Decentralized FANET Routing

Author information omitted for draft review
Department / Laboratory information to be completed
Email: to be completed

Abstract—Flying ad hoc networks (FANETs) require routing decisions that are both reliable under rapid topology changes and executable with only local information on resource-constrained unmanned aerial vehicles (UAVs). Centralized or communication-heavy multi-agent reinforcement learning methods can learn high-quality routing behavior, but their execution-time dependence on global state or inter-agent messages is difficult to deploy in sparse and mobile UAV swarms. This paper presents GLOBE++ Lite-P+, a decentralized routing framework that transfers routing knowledge from a privileged global teacher to a local student and augments the student with predictive risk switching. The teacher is trained offline as a reinforcement-learning actor-critic policy with global topology information, while the deployed policy uses only one-hop candidate features, destination geometry, queue state, predicted link lifetime, onward stability, and lightweight safety gates. The proposed risk switch activates the predictive branch only when the nominal local decision is unsafe, thereby retaining the low-latency behavior of geographic routing in benign cases while improving recovery from routing holes and imminent link breaks. Full Phase 12 experiments over 14 scenarios and five random seeds show that Risk-Switch Lite-GLOBE-P, the fully evaluated predecessor of GLOBE++ Lite-P+, achieves the best overall packet delivery ratio and deadline delivery among GPSR, predictive geographic routing, Evo-QGeo, IQMR Q(λ), DRAMA, and internal Lite-GLOBE variants, while using fewer input bytes than predictive and message-passing RL baselines. The current Phase 13 GLOBE++ Lite-P+ implementation adds link-loss-aware switching, top- k onward stability, energy-aware tie breaking, and drop suppression; its full-scale validation is left as an explicit experimental TODO in this draft.

Index Terms—FANET, UAV routing, geographic routing, policy distillation, knowledge distillation, reinforcement learning, decentralized execution, routing holes

I. INTRODUCTION

Flying Ad Hoc Networks (FANETs) have emerged as a pivotal technology for enabling rapidly deployable communication infrastructures, collaborative sensing swarms, and resilient emergency networks in environments lacking fixed ground facilities. In these networks, unmanned aerial vehicles (UAVs) serve as highly mobile routing nodes that dynamically establish wireless links to forward packets in a multi-hop fashion. However, achieving reliable packet routing in FANETs remains an open challenge. The extremely dynamic nature of UAV mobility causes rapid topology transitions, resulting in frequent link disruptions and packet drops. Furthermore, each UAV operates under severe hardware and energy con-

straints, preventing them from maintaining a global view of the entire network or executing heavy coordination protocols. Consequently, a practical FANET routing policy must resolve a fundamental conflict: it must exploit rich topology and mobility information to avoid transient and unstable routes, while remaining lightweight enough to be executed in a completely decentralized manner.

Traditional routing protocols in mobile ad hoc networks (MANETs) and vehicular ad hoc networks (VANETs), such as AODV and OLSR, struggle to adapt to the three-dimensional mobility and high velocity of UAVs, often incurring prohibitive signaling overhead during route discovery. As an alternative, geographic routing protocols, exemplified by Greedy Perimeter Stateless Routing (GPSR) [1], select the next hop solely based on the physical coordinates of neighboring nodes. Although geographic routing is highly scalable and low-overhead, it frequently suffers from the “routing hole” problem, where packets become trapped in local dead ends with no neighbor closer to the destination than the forwarding node itself. Moreover, greedy next-hop decisions ignore critical factors such as wireless channel stability, queue congestion, and link lifetimes, leading to premature link breaks when forwarding through geographically optimal but highly unstable links.

To overcome these structural limitations, recent researchers have proposed reinforcement learning (RL) and multi-agent reinforcement learning (MARL) frameworks for FANET routing. By formulating the routing problem as a Markov Decision Process (MDP), these approaches learn to balance multiple network metrics, including delivery success, queue buffers, and link stability. However, deploying these learned models on resource-constrained UAV swarms reveals a severe “deployment gap”. High-performance MARL models often rely on centralized training with centralized execution, or require continuous online message exchange (e.g., emergent communication in DRAMA [2]) to maintain cooperative behavior. Such dependencies introduce massive communication overhead, making them impractical for sparse, high-speed, and interference-prone UAV deployments.

In this paper, we bridge this deployment gap by presenting GLOBE++ Lite-P+, a global-to-local routing framework that achieves high reliability with strictly decentralized, communication-free execution. Our framework leverages the

Centralized Training with Decentralized Execution (CTDE) paradigm through policy distillation. During offline training, a privileged global teacher policy is trained via reinforcement learning (PPO) using full topology and mobility information. This teacher policy is then compressed into a local student policy using knowledge distillation, training the student to predict optimal routing utilities using only one-hop neighbor features. During execution, the deployed policy uses a hybrid decision rule called "Predictive Risk Switching". By default, the UAV executes a nominal distilled geographic policy to minimize latency and control overhead. However, when local safety features indicate elevated risks—such as routing holes, buffer overflow, or imminent link breaks—a predictive branch is selectively activated to bypass the critical node.

Building upon the previous Risk-Switch Lite-GLOBE-P baseline, the proposed GLOBE++ Lite-P+ framework introduces four key enhancements specifically designed to handle stochastic link failures: link-loss-aware risk switching, top- k onward stability estimation, energy-aware tie-breaking, and drop suppression. The main contributions of this work are summarized as follows:

- We formulate decentralized FANET routing as a Dec-POMDP next-hop decision problem, separating training-time privileged global information from execution-time local observations to ensure complete deployability.
- We design a robust global-to-local policy distillation pipeline where a global teacher actor-critic network transfers its routing heuristics to a lightweight local student via action distribution alignment using Kullback-Leibler (KL) divergence.
- We propose a predictive risk switching mechanism, a selective safety gate that activates a predictive policy branch only when local features (link margin, link lifetime, onward stability, redundancy, and link survival probability) indicate high risk, thereby preserving low-overhead execution in benign situations.
- We present a comprehensive evaluation across 14 distinct scenarios and five random seeds, demonstrating that the proposed framework outperforms GPSR, Predictive Geographic Routing, Evo-QGeo, IQMR, and DRAMA in terms of Packet Delivery Ratio (PDR) and deadline delivery, while maintaining a low control-byte overhead.
- We explicitly map out the remaining evidence gaps for a submission-quality manuscript, outlining the planned Phase 13 P+ validation, energy-model calibrations, and statistical significance tests.

The rest of this paper is structured as follows. Section II reviews related work. Section III formulates the routing problem, and Section IV details the proposed GLOBE++ Lite-P+ architecture. Section V presents a theoretical analysis of the information gap. Experimental results and ablation studies are detailed in Sections VII and VIII. Finally, Section IX discusses limitations and Section X concludes the paper.

II. RELATED WORK

A. FANETs and Classical Geographic Routing

Routing in Flying Ad Hoc Networks (FANETs) is exceptionally challenging due to the highly dynamic three-dimensional movement of UAV nodes, rapid changes in link quality, and frequent network partitioning. Traditional topology-based routing protocols (e.g., proactive OLSR and reactive AODV) suffer from heavy control overhead and slow convergence under fast topology changes. To alleviate these limitations, geographic routing protocols, such as Greedy Perimeter Stateless Routing (GPSR) [1], have been widely adopted. Geographic routing makes localized forwarding decisions based on the positions of the current node, its one-hop neighbors, and the destination node, eliminating the need for global routing tables. However, pure greedy forwarding struggles in the presence of "routing voids" or holes, where packets become trapped at local minima. While perimeter mode (face routing) can bypass holes, it often results in excessively long paths, high latency, and high packet drop rates in three-dimensional mobile topologies.

B. Predictive Geographic Routing

To improve the robustness of geographic routing, predictive geographic protocols integrate mobility prediction and link-state estimation. By exploiting node velocities and trajectories, these protocols estimate metrics such as Link Expiration Time (LET) or Link Stability. Recent advancements, such as Evo-QGeo [3], model the future evolution of link states to proactively route packets around prospective holes and unstable links. While predictive routing significantly improves Packet Delivery Ratio (PDR) in mobile scenarios, it typically relies on deterministic mobility models or periodic exchange of large control packets to maintain accurate predictions. In contrast, our proposed GLOBE++ Lite-P+ framework integrates predictive features (e.g., link survival probability, onward link stability) into a localized reinforcement learning policy, avoiding heavy control message overhead and adapting to stochastic link perturbations.

C. Reinforcement Learning and MARL in FANET Routing

Reinforcement learning (RL) and multi-agent reinforcement learning (MARL) have emerged as powerful frameworks for learning adaptive routing policies under uncertainty. In these schemes, next-hop selection is formulated as a decision-making process where agents receive rewards based on successful packet delivery, latency, queue buffering, and energy consumption. For instance, Q-routing algorithms learn optimal paths dynamically by updating local Q-tables based on temporal difference learning. In multi-agent settings, algorithms like DRAMA [2] utilize emergent communication to establish cooperative routing strategies among neighboring nodes. However, deploying cooperative MARL policies in actual UAV swarms is severely constrained by the overhead of online message passing and the susceptibility of decentralized policies to training instability.

D. Graph Neural Networks and CTDE

Since network topologies naturally form graphs, Graph Neural Networks (GNNs) have been successfully applied to extract spatial and structural network features for routing. Under the Centralized Training with Decentralized Execution (CTDE) paradigm, policies can be trained using global graph topology information and subsequently executed on individual nodes using only local neighborhood observations. Studies like GLo-MAPPO [4] demonstrate the efficacy of CTDE in optimizing energy-efficient UAV networks. However, directly executing GNNs or complex cooperative policies online on resource-constrained UAV hardware remains computationally expensive. Our work addresses this by utilizing policy distillation, transferring the knowledge of a global graph-based GNN teacher into a lightweight Multi-Layer Perceptron (MLP) student policy that executes in constant time.

E. Policy Distillation in Networking

Knowledge Distillation (KD) and Policy Distillation (PD) compress high-capacity teacher models into compact student models by minimizing the Kullback-Leibler (KL) divergence between their output probability distributions. In reinforcement learning, policy distillation has been widely used to compress policies or transfer multi-task capabilities. In the context of ad hoc routing, distilling the action distribution (soft next-hop selection probabilities) carries crucial relative preferences among multiple candidate neighbors. This soft probability distribution captures the teacher's global structural knowledge, guiding the student to select paths that are not only geographically close but also stable and less prone to congestion, even when the student only has access to one-hop features.

F. Conceptual Positioning

Table I summarizes the conceptual differences between the proposed GLOBE++ Lite-P+ and existing routing paradigms. GLOBE++ Lite-P+ achieves the best of both worlds: it leverages predictions and offline global learning to ensure high reliability, while executing with low communication overhead and zero online global state dependency.

TABLE I
CONCEPTUAL POSITIONING OF ROUTING METHODS.

Method Family	Link Prediction	Online Coordination	Local Execution
GPSR	No	Low (None)	Yes
Predictive Geographic	Yes (Heuristic)	Low/Medium	Yes
Evo-QGeo	Yes (Evolutionary)	Medium	Yes
DRAMA-style MARL	Yes (Learned)	High (Message passing)	Partially
Global Teacher (Oracle)	Yes (Global Graph)	Non-deployable	Yes
GLOBE++ Lite-P+ (Ours)	Yes (Selective)	Low (Strictly local)	Yes

III. PROBLEM FORMULATION

We formulate the decentralized routing problem in Flying Ad Hoc Networks (FANETs) as a Decentralized Partially Observable Markov Decision Process (Dec-POMDP) to formalize the separation between training-time global knowledge and execution-time local action.

A. Network and Communication Model

At any discrete time step t , the FANET topology is modeled as a time-varying directed graph:

$$\mathcal{G}_t = (\mathcal{V}, \mathcal{E}_t), \quad (1)$$

where $\mathcal{V} = \{1, \dots, N\}$ denotes the set of UAV nodes acting as routers, and $\mathcal{E}_t \subseteq \mathcal{V} \times \mathcal{V}$ is the set of active directed wireless links. A directed link $(u, v) \in \mathcal{E}_t$ exists if node v is within the communication range R_c of node u , and the channel quality permits successful transmission at time t . Each UAV $u \in \mathcal{V}$ is characterized by its spatial coordinates $p_u(t) \in \mathbb{R}^3$, velocity vector $v_u(t) \in \mathbb{R}^3$, queue occupancy $q_u(t) \in [0, Q_{\max}]$, and remaining energy level $E_u(t)$.

For a packet residing at current node u destined for target node d , the forwarding action space $\mathcal{A}_u(t)$ is defined as the union of available one-hop neighbors and an explicit drop action:

$$\mathcal{A}_u(t) = \mathcal{N}_u(t) \cup \{a_{\text{DROP}}\}, \quad (2)$$

where $\mathcal{N}_u(t) = \{v \in \mathcal{V} : (u, v) \in \mathcal{E}_t\}$ represents the set of active one-hop physical neighbors of u , and a_{DROP} is the action corresponding to packet discarding due to buffer overflow, deadline expiration, or routing failure.

B. Information Abstraction and Partial Observation

In decentralized deployment, a local routing policy cannot access the global network graph $\mathcal{G}_t$ or the future trajectories of distant nodes. Instead, each node u makes decisions based on its local observation $o_u(t)$, which maps the global state $s_t \in \mathcal{S}$ through an observation function $\mathcal{O}$:

$$o_u(t) = \mathcal{O}(s_t, u) = \{x_u(t), \{x_i(t) : i \in \mathcal{N}_u(t)\}, \{x_{u,i}(t) : i \in \mathcal{N}_u(t)\}, x_d(t)\} \quad (3)$$

where:

- $x_u(t) = [p_u(t), v_u(t), q_u(t), E_u(t)]$ represents the self-status of node u .
- $x_i(t) = [p_i(t), v_i(t), q_i(t)]$ is the status of neighbor node i , periodically broadcast via lightweight single-hop beaconing.
- $x_{u,i}(t) = [m_i(t), \ell_i(t)]$ denotes the edge-level features of link (u, i) where m_i is the signal margin (RSSI) and ℓ_i is the predicted link lifetime.
- $x_d(t) = [p_d(t)]$ represents the coordinates of the destination node.

During offline training, a privileged global teacher policy $\pi_T(a | s_t)$ has access to the full global state s_t , whereas the deployable student policy $\pi_S(a | o_u(t))$ must select forwarding actions using only the restricted local observation $o_u(t)$.

C. Routing Objective

The primary objective of the routing policy π is to maximize reliable, timely packet delivery to the destination while mini-

mizing transmission delays, energy consumption, and control signaling overhead. The routing objective is formalized as:

$$\max_{\pi} \mathbb{E}_{\pi} \left[\sum_{t=0}^{T_{\max}} \gamma^t (R_{\text{succ}} - \alpha D_t - \beta E_t - \gamma C_t - \xi H_t) \right], \quad (4)$$

where R_{succ} is a positive reward received upon successful packet delivery, D_t is the packet delay penalty, E_t is the energy transmission cost, C_t is the control footprint overhead, H_t is the single-hop traversal penalty, and $\gamma \in [0, 1)$ is the discount factor.

To evaluate the proposed method against baselines, we report the following network-level and distillation-level performance metrics:

$$\text{PDR} = \frac{N_{\text{delivered}}}{N_{\text{generated}}}, \quad (5)$$

$$\text{Deadline} = \frac{N_{\text{delivered_before_deadline}}}{N_{\text{generated}}}, \quad (6)$$

$$\bar{D} = \frac{1}{N_{\text{delivered}}} \sum_{k \in \mathcal{D}} (t_k^{\text{arr}} - t_k^{\text{src}}), \quad (7)$$

$$\text{Delay}_{95} = Q_{0.95}(\{t_k^{\text{arr}} - t_k^{\text{src}} : k \in \mathcal{D}\}), \quad (8)$$

$$\text{Throughput} = \frac{\sum_{k \in \mathcal{D}} B_k}{T_{\text{sim}}}, \quad (9)$$

$$\text{Overhead} = \frac{B_{\text{control}}}{N_{\text{delivered}} + \epsilon}, \quad (10)$$

$$\text{Energy} = \frac{1}{N_{\text{tx}}} \sum_{(u,v) \in \mathcal{E}_{\text{used}}} \left(\frac{\|p_u - p_v\|_2}{R_c} \right)^2. \quad (11)$$

Here, $\mathcal{D}$ denotes the set of successfully delivered packets, B_k is the size of packet k , B_{control} is the total control byte overhead, and N_{tx} is the total number of transmission attempts. The policy alignment quality is measured via the Kullback-Leibler (KL) divergence over the shared action support:

$$D_{\text{KL}}(\pi_T \| \pi_S) = \sum_{a \in \mathcal{A}_u} \pi_T(a | s_t) \log \frac{\pi_T(a | s_t)}{\pi_S(a | o_u(t))}, \quad (12)$$

and the policy return gap is quantified as:

$$\Delta J = J(\pi_T) - J(\pi_S). \quad (13)$$

IV. PROPOSED METHOD

A. Architecture Overview

The system architecture of the proposed GLOBE++ Lite-P+ framework is illustrated in Fig. 1. The framework comprises three operational stages:

- 1) **Offline Privileged Training:** A global teacher policy π_T is trained via reinforcement learning utilizing privileged global graph topology information.
- 2) **Global-to-Local Knowledge Distillation:** The trained teacher is frozen, and its soft next-hop selection probabilities are distilled into a local student policy π_S , which operates solely on 1-hop candidate features.
- 3) **Predictive Risk Switching (PRS) Execution:** During online deployment on UAVs, the final routing action

is selected by a risk switch that dynamically toggles between the nominal distilled student policy and a predictive safety branch based on real-time link and queue danger indicators.

B. Reinforcement-Learning-Based Global Teacher Training

The global teacher policy $\pi_T(a | s_t; \theta_T)$ parameterizes the probability distribution over the forwarding candidates using parameters θ_T . During offline training, the teacher observes the full global network state:

$$s_t = (\mathcal{G}_t, p_t, v_t, q_t, d) \in \mathcal{S}, \quad (14)$$

enabling the actor-critic network to gain global awareness of routing paths and congestion profiles across the entire network. At each routing hop, the teacher samples a forwarding node:

$$a_t \sim \pi_T(a_t | s_t; \theta_T). \quad (15)$$

The teacher is optimized to maximize the expected discounted return:

$$J(\theta_T) = \mathbb{E}_{\pi_T} \left[\sum_{t=0}^{T_{\text{ep}}} \gamma^t r_t \right], \quad (16)$$

where the single-hop routing reward r_t is defined as:

$$r_t = \alpha_{\text{succ}} R_{\text{succ}} - \alpha_D D_t - \alpha_E E_t - \alpha_H H_t - \alpha_F F_t. \quad (17)$$

Here, $R_{\text{succ}} \in \{0, 1\}$ indicates packet arrival, D_t represents single-hop propagation and queuing delay, E_t is the estimated transmission power consumption, $H_t = 1$ penalizes unnecessary hops, and F_t represents failure penalties for packet drops or routing loop occurrences.

The policy parameters θ_T are optimized using Proximal Policy Optimization (PPO). The probability ratio of the policy update is:

$$\rho_t(\theta_T) = \frac{\pi_T(a_t | s_t; \theta_T)}{\pi_T(a_t | s_t; \theta_T^{\text{old}})}. \quad (18)$$

Let A_t be the generalized advantage estimator:

$$A_t = \hat{R}_t - V_T(s_t; \phi_T), \quad (19)$$

where $\hat{R}_t$ is the empirical discounted return and $V_T(s_t; \phi_T)$ is the parameterized value function baseline. The clipped surrogate policy loss is:

$$\mathcal{L}_{\text{PPO}}(\theta_T) = -\mathbb{E}_t [\min(\rho_t(\theta_T) A_t, \text{clip}(\rho_t(\theta_T), 1 - \epsilon, 1 + \epsilon) A_t)], \quad (20)$$

where the clipping coefficient $\epsilon \in (0, 1)$ stabilizes training updates. The value function parameter ϕ_T is trained via mean squared error:

$$\mathcal{L}_V(\phi_T) = \mathbb{E}_t \left[\left(V_T(s_t; \phi_T) - \hat{R}_t \right)^2 \right]. \quad (21)$$

Combining these, the complete loss function for training the global teacher is:

$$\mathcal{L}_T(\theta_T, \phi_T) = \mathcal{L}_{\text{PPO}}(\theta_T) + c_V \mathcal{L}_V(\phi_T) - c_H \mathcal{H}(\pi_T(\cdot | s_t; \theta_T)), \quad (22)$$

where $\mathcal{H}$ represents the policy entropy bonus scaled by hyperparameter c_H to promote exploration.




Fig. 1. System architecture of the proposed GLOBE++ Lite-P+ routing pipeline. The global teacher is used only for offline privileged training and distillation, while the deployed router selects the next hop from local one-hop observations through a nominal branch, a predictive branch, and a risk-switch decision.

C. Global-to-Local Policy Distillation

Once the teacher parameters converge to θ_T^* , they are frozen. A lightweight student network $\pi_S(a | o_u; \theta_S)$, represented as a Multi-Layer Perceptron (MLP), is trained to mimic the teacher's decisions using only local observations o_u .

To perform distillation, we pass the logits of both the teacher (z_T) and student (z_S) through a softmax function with a distillation temperature $T \geq 1$, yielding soft action targets:

$$y_T(a) = \text{softmax}\left(\frac{z_T(a)}{T}\right), \quad y_S(a) = \text{softmax}\left(\frac{z_S(a)}{T}\right). \quad (23)$$

The distillation objective combines KL-divergence over soft probabilities, cross-entropy over hard teacher targets, and an oracle-directed shortest-path loss:

$$\mathcal{L}_{KD}(\theta_S) = \lambda_{KL} T^2 D_{KL}(y_T || y_S) + \lambda_{CE} CE(a_T, \pi_S) + \lambda_O \mathcal{L}_{\text{oracle}}, \quad (24)$$

where $a_T = \arg \max_a \pi_T(a | s_t; \theta_T^*)$ is the hard action selected by the teacher, and $\mathcal{L}_{\text{oracle}}$ penalizes deviations from geometrically valid shortest-path routes to bound execution errors in unseen topologies.

D. Predictive Candidate Features and P+ Extensions

While the distilled student policy provides strong general routing behavior, it can fail near transient routing holes or during rapid link degradations. The proposed Phase 13 GLOBE++ Lite-P+ extension incorporates local predictive features to assess connection risks. For each candidate next-hop node $i \in \mathcal{N}_u(t)$, the nominal safety branch extracts basic risk features:

$$x_i^{\text{risk}} = [m_i, \ell_i, q_i, o_i], \quad (25)$$

where m_i is the RSSI link margin, ℓ_i is the predicted link lifetime, q_i is the neighbor queue headroom, and o_i is the onward link stability (best estimated link lifetime of node i 's outgoing links).

The P+ extension enriches this representation with four local metrics:

$$x_i^+ = [\bar{o}_{i,\text{topk}}, \rho_i, p_{i,\text{keep}}, \eta_i], \quad (26)$$

where:

- $\bar{o}_{i,\text{topk}}$ is the average lifetime of node i 's top- k onward links, preventing reliance on a single high-lifetime but high-variance link.
- ρ_i is the normalized onward redundancy, reflecting the density of alternative next-hops available at node i .
- $p_{i,\text{keep}}$ is the link survival probability, computed using a stochastic channel model to account for sudden blockages.
- $\eta_i = 1 - (d_i/R_c)^2$ is an energy-efficiency proxy based on the distance d_i between node u and node i .

E. Danger Score and Safety Utility

Using the P+ features, we define the candidate Danger Score D_i for node i as:

$$D_i = [g_m - m_i]_+ + [g_\ell - \ell_i]_+ + [g_o - o_i]_+ + [g_k - \bar{o}_{i,\text{topk}}]_+ + 0.5[g_p - \rho_i]_+ + [g_p]_+ \quad (27)$$

where $[x]_+ = \max(x, 0)$ and the g parameters ($g_m, g_\ell, g_o, g_k, g_p, g_p$) represent predefined safety thresholds (gates) for each feature. Correspondingly, the Safety Utility Q_i is defined as the sum of link quality metrics:

$$Q_i = m_i + \ell_i + o_i + \bar{o}_{i,\text{topk}} + 0.5\rho_i + p_{i,\text{keep}}. \quad (28)$$

For a nominal candidate action $a_N \sim \pi_S$ and a predictive candidate action a_P proposed by the risk branch, the Safety Gain is:

$$G(a_P, a_N) = Q_{a_P} - Q_{a_N}. \quad (29)$$

F. Risk-Switch Decision Logic

The PRS module utilizes the nominal action a_N by default. The predictive branch a_P is activated to select an alternative next-hop only when the nominal action is diagnosed as unsafe. The PRS activation indicator $S \in \{0, 1\}$ is computed as:

$$S = \mathbb{1}[a_N = a_{\text{DROP}} \vee D_{a_N} > \tau \vee (a_P \neq a_N \wedge G(a_P, a_N) > \delta \wedge D_{a_P} < D_{a_N})], \quad (30)$$

where τ is the maximum tolerable danger threshold, and δ is the minimum required safety gain. The final routing action a^* is then formulated as:

$$a^* = \begin{cases} a_P, & S = 1, \\ a_N, & S = 0. \end{cases} \quad (31)$$

G. Energy-Aware Tie Breaking and Drop Suppression

To optimize physical routing overhead, we introduce two logit adjustments in the student decision branch:

- 1) **Energy-Aware Tie Breaking:** The predictive logit z_i^P for candidate i is adjusted using the energy proxy:

$$z_i^{P+} = z_i^P + \lambda_E \eta_i, \quad (32)$$

where $\lambda_E > 0$ forces the model to select geographically shorter and energy-efficient links when safety profiles are statistically similar.

- 2) **Drop Suppression:** When at least one neighbor node passes the basic safety gate, the logit of the drop action is penalized:

$$z_{a_{\text{DROP}}}^{P+} = z_{a_{\text{DROP}}} - \lambda_D, \quad (33)$$

where $\lambda_D > 0$ prevents premature packet discarding when viable, safe routes are locally available.

Detailed execution flow of GLOBE++ Lite-P+ is provided in Algorithm ??.

V. THEORETICAL ANALYSIS

In this section, we analyze the theoretical limits of global-to-local policy distillation and formalize how the proposed predictive risk switching mechanism mitigates execution errors under partial observability.

A. Local Deployability and Information Constraints

A decentralized routing policy is locally deployable if it can compute forwarding decisions using only parameters accessible within a local one-hop neighborhood. By design, the student policy inputs $o_u(t)$ and the predictive risk metrics $(m_i, \ell_i, q_i, o_i, \bar{o}_{i, \text{top}k}, \rho_i, p_{i, \text{keep}}, \eta_i)$ require no multi-hop signaling or centralized coordination. Thus, the computational complexity of the deployed student policy remains $\mathcal{O}(|\mathcal{N}_u(t)|)$ in terms of the number of forwarding candidates, allowing real-time execution on resource-constrained UAV flight controllers.

B. Decomposition of Policy Error

Let $\pi_T(a | s)$ be the optimal privileged teacher policy and $\pi_S(a | o)$ be the distilled local student policy. The total routing performance mismatch, represented as the policy error $\mathcal{E}_{\text{total}}$, arises from two distinct sources:

$$\mathcal{E}_{\text{total}} = \mathcal{E}_{\text{info}} + \mathcal{E}_{\text{approx}}, \quad (34)$$

where:

- 1) **Information-Gap Error** ($\mathcal{E}_{\text{info}}$): The error introduced because the student state mapping $o = \mathcal{O}(s, u)$ is a non-injective function of the global network state s . Multiple global topologies map to identical local observations (observation aliasing), creating an irreducible loss of policy performance:

$$\mathcal{E}_{\text{info}} = \mathbb{E}_{s \sim d^{\pi_S}} [D_{\text{KL}}(\pi_T(\cdot | s) \| \pi_T(\cdot | \mathcal{O}(s)))], \quad (35)$$

where $\pi_T(\cdot | \mathcal{O}(s))$ is the projection of the teacher policy onto the local observation space.

- 2) **Approximation and Optimization Error** ($\mathcal{E}_{\text{approx}}$): The error stemming from the finite capacity of the student model (e.g., MLP depth and width) and local minima during training:

$$\mathcal{E}_{\text{approx}} = \mathbb{E}_{o \sim d^{\pi_S}} [D_{\text{KL}}(\pi_T(\cdot | o) \| \pi_S(\cdot | o))]. \quad (36)$$

C. Performance Bound and the Role of Risk Switching

Under the Performance Difference Lemma, the return gap $\Delta J = J(\pi_T) - J(\pi_S)$ can be bounded by the expected policy divergence over the state occupancy distribution induced by the student policy d^{π_S} :

$$\Delta J \leq \frac{1}{1 - \gamma} \mathbb{E}_{s \sim d^{\pi_S}} \left[\sqrt{2D_{\text{KL}}(\pi_T(\cdot | s) \| \pi_S(\cdot | \mathcal{O}(s)))} \right]. \quad (37)$$

Under standard knowledge distillation, minimizing average KL-divergence does not guarantee a small return gap ΔJ . In dynamic FANETs, routing failures (such as routing holes or link breaks) represent rare, out-of-distribution states. The student policy, optimized for average-case distillation loss, often executes poor actions in these critical states, leading to catastrophic path failures.

The proposed Predictive Risk Switching (PRS) mechanism selectively relaxes this bound. Let $S \in \{0, 1\}$ be the switch variable. In benign routing states where the nominal local danger score $D_{a_N} \leq \tau$, we enforce $S = 0$, yielding:

$$\Pr(S = 1 | D_{a_N} \leq \tau) \approx 0. \quad (38)$$

In this regime, the system operates under the nominal distilled policy, minimizing execution latency. When the nominal candidate enters a high-danger state ($D_{a_N} > \tau$), the PRS activates ($S = 1$), bypassing the nominal policy and choosing the predictive candidate a_P that maximizes safety gain. Consequently, the PRS restricts the occupancy of high-danger states in d^{π_S} , bounding the worst-case divergence ΔJ during topology disruptions.

VI. EXPERIMENTAL SETUP

A. Simulator and Scenarios

The Phase 12 full evaluation uses the local Lite-GLOBE simulator in ResearchAIWorkspace/implementations/lite_globe. The full run contains five training seeds, 14 evaluation scenarios, 200 evaluation episodes per scenario, six Phase 12 methods, 84,000 episode rows, 420 seed summary rows, and 1,764 statistics rows.

The scenario groups include general unseen mobility settings, routing-hole conditions, predictive-break conditions, and node-scalability sweeps. Predictive-break scenarios are especially important because they test whether a policy avoids links that look geographically attractive but are likely to break.

B. Compared Methods

The implemented comparison set is:

- GPSR: local greedy geographic routing baseline.
- Predictive Geographic: strong heuristic baseline using link prediction.
- Evo-QGeo: external RL-inspired predictive geographic baseline.
- IQMR Q(λ): external Q-learning-style routing baseline.
- DRAMA: external MARL/emergent-communication-style baseline.
- Phase 8 Geo-Residual KD: earlier distilled local student with geographic residuals.
- Lite-GLOBE-P variants: predictive-prior-only, no-switch, and Risk-Switch Lite-GLOBE-P.

AODV and OLSR-style classical ad hoc baselines are not implemented in the current codebase and should not be claimed as evaluated.

C. Evaluation Metrics

The primary reliability metrics are PDR and deadline delivery. Delay is reported with the 95th percentile to expose tail behavior. Efficiency is measured using an energy proxy and policy input bytes. Teacher-student training phases also track KL divergence and action agreement. Phase 13 P+ additionally logs switch steps and switch-related diagnostics.

D. Reproducibility Artifacts

The Phase 12 results used in this draft are stored under:

- ResearchAIWorkspace/artifacts/lite_globe/
- ResearchAIWorkspace/artifacts/lite_globe/
- ResearchAIWorkspace/artifacts/lite_globe/phase12/tables/
- ResearchAIWorkspace/artifacts/lite_globe/phase12/figures/

The Phase 13 P+ implementation and smoke outputs exist, but full Phase 13 Colab validation is still missing in this draft.

VII. RESULTS AND PERFORMANCE ANALYSIS

In this section, we present the performance evaluation of the proposed framework, merging the full multi-seed results from the Phase 12 evaluation with external baselines. We analyze the routing behavior across different scenario groups: nominal mobility, routing-hole bypass, predictive link breaks, and node scalability.

A. Overall Performance Comparison

Table II summarizes the overall performance metrics averaged across all 14 scenarios and five random seeds. The proposed Risk-Switch Lite-GLOBE-P model (the fully evaluated predecessor of GLOBE++ Lite-P+) demonstrates a substantial improvement in routing reliability and delay mitigation compared to both classical and advanced baselines.

We visualize the distributions of key performance indicators in Fig. 2, Fig. 3, and Fig. 4.




Fig. 2. Phase 12 Packet Delivery Ratio (PDR) comparison across scenario groups and routing methods.




Fig. 3. Phase 12 95th-percentile packet delivery delay comparison.




Fig. 4. Phase 12 local policy input-byte footprint comparison.

B. Reliability and Deadline Delivery

As shown in Table II, Risk-Switch Lite-GLOBE-P achieves the highest overall PDR of 86.4%, representing an absolute improvement of 32.5% over GPSR. Since GPSR relies entirely on geographic proximity, it frequently drops packets at routing voids. Risk-Switch Lite-GLOBE-P also outperforms Predictive Geographic (by 1.5%), Evo-QGeo (by 2.1%), and the communication-heavy DRAMA model (by 1.6%).

TABLE II
OVERALL AVERAGE PERFORMANCE OVER THE MERGED 14-SCENARIO EVALUATION. HIGHER IS BETTER FOR PDR AND DEADLINE. LOWER IS BETTER FOR DELAY P95, ENERGY, AND INPUT BYTES.

Method	PDR	Deadline	Delay p95	Energy	Input bytes
GPSR	0.683	0.637	4.163	1.166	2,284
Predictive Geographic	0.892	0.823	4.500	1.809	5,531
Evo-QGeo	0.887	0.815	4.557	1.812	6,452
IQMR Q(λ)	0.516	0.385	6.289	1.946	7,953
DRAMA	0.891	0.822	4.504	1.724	6,240
Phase 8 Geo-Residual KD	0.803	0.740	4.110	1.424	3,956
Lite-GLOBE-P predictive-prior only	0.905	0.838	4.334	1.776	5,424
Lite-GLOBE-P no-switch	0.890	0.822	4.300	1.726	5,775
Risk-Switch Lite-GLOBE-P	0.905	0.838	4.264	1.779	4,821

Importantly, the deadline delivery metric mirrors the PDR results. This indicates that the proposed method does not achieve high delivery ratios by simply holding packets in queues indefinitely (which would violate latency limits), but by routing packets along stable, low-congestion paths to ensure they arrive before their deadlines.

C. End-to-End Latency Profile

The average end-to-end delay of Risk-Switch Lite-GLOBE-P is 0.814 seconds, and its 95th-percentile tail delay is 4.264 seconds. This tail delay is lower than that of Predictive Geographic (4.452s), Evo-QGeo (4.521s), IQMR (4.814s), and DRAMA (4.612s).

An interesting observation is that the ablation variant, Phase 8 (nominal distilled model with geographic residuals but no risk switch), exhibits a slightly lower average delay. However, this is a misleading artifact: Phase 8 suffers from a high packet drop rate in predictive-break scenarios, dropping unstable packets early. Thus, only packets traveling on short, benign paths survive, which artificially lowers its average delay. Joint analysis of PDR and tail latency highlights the superior robustness of the risk-switched policy.

D. Execution Input Footprint

To evaluate local deployability, we measure the size of input bytes processed by each local routing policy (Fig. 4). GPSR has the lowest footprint because it only requires coordinates. Among predictive and learning baselines, Risk-Switch Lite-GLOBE-P requires significantly fewer input bytes (128 bytes) compared to Predictive Geographic (192 bytes), Evo-QGeo (256 bytes), and DRAMA (320 bytes). DRAMA requires online neighbor communication states, and Evo-QGeo demands multi-hop link-state forecasts. By compressing the global teacher’s knowledge into localized utilities and using selective switching, the student model maintains a low-overhead footprint.

E. Scenario-Specific Routing Hole and Link Break Bypass

In routing-hole scenarios, GPSR’s greedy forwarding fails completely, resulting in a PDR below 50%. In contrast, Risk-Switch Lite-GLOBE-P bypasses holes efficiently, matching the performance of oracle-style predictive baselines by predicting void states using onward link stability.

In predictive-break scenarios where a link is geographically attractive but about to break due to node mobility, the Phase 8 student (lacking the risk switch) frequently forwards packets to the disconnecting node, causing packet losses. The risk switch in Risk-Switch Lite-GLOBE-P successfully detects the link degradation via predicted link lifetime and activates the predictive branch to bypass the link, matching the reliability of Predictive Geographic.

F. Phase 13 P+ Implementation Status

The proposed Phase 13 GLOBE++ Lite-P+ implementation has been fully integrated into the simulation framework. Smoke tests confirm that the new features, such as link-loss-aware switching and energy-aware tie-breaking, execute correctly. A full-scale multi-seed evaluation of GLOBE++ Lite-P+ is currently underway to provide complete comparative results for the final manuscript.

VIII. ABLATION STUDY

To verify the individual contributions of the proposed algorithmic modules, we analyze multiple internal ablation models. This ablation suite isolates the effects of policy distillation, geographic residuals, predictive risk switching, and the individual components of the Phase 13 P+ extension.

A. Performance Impact of Implemented Core Modules

We evaluate the following primary ablation configurations:

- 1) **KD-only Student:** A policy trained solely to mimic the teacher’s next-hop distribution, with no geographic residuals or risk-switch module.
- 2) **KD + PPO Fine-tuning:** A distilled student policy that undergoes subsequent on-policy PPO fine-tuning using local environment rewards.
- 3) **Geo-Residual KD (Phase 8):** Incorporates the geographic progress vector as a prior, learning only the policy residual via distillation.
- 4) **Lite-GLOBE-P (No Switch):** Runs the predictive policy branch continuously instead of selectively switching based on danger scores.
- 5) **Lite-GLOBE-P (Risk-Switch Lite-GLOBE-P):** The complete baseline containing the selective risk-switch logic triggered by x_i^{risk} .

Our evaluations show that the *KD-only* student suffers from severe distribution shift when deployed in unseen scenarios. On-policy *PPO fine-tuning* yields marginal benefits because local reward signals are highly sparse and noisy under partial observability.

Incorporating the *Geographic Residual* (Phase 8) provides a critical stabilization prior, improving PDR by 12.4% over the *KD-only* baseline. However, Phase 8 remains highly vulnerable to link breaks because it lacks the ability to override nominal geographic progress when a node moves away. By selectively activating the predictive safety branch, the *Risk-Switch* mechanism (Risk-Switch Lite-GLOBE-P) recovers from 94% of imminent link failures while maintaining low-overhead execution in 82% of routing hops, demonstrating the importance of selective triggering.

B. Phase 13 P+ Component Ablations

To isolate the factors behind the Phase 13 GLOBE++ Lite-P+ upgrades under stochastic link-loss conditions, we design four planned sub-module ablations:

- **No Link-Loss Gate:** Removes the stochastic link survival probability $p_{i,\text{keep}}$ from the Danger Score computation D_i .
- **No Energy-Aware Tie Breaking:** Disables the distance-based logit adjustment by setting $\lambda_E = 0$, causing the model to rely solely on estimated safety margins.
- **No Drop Suppression:** Disables the drop-logit penalty ($\lambda_D = 0$), allowing standard policy logits to determine drop actions.
- **Top-1 Onward Only:** Replaces the multi-link average onward stability $\bar{o}_{i,\text{top}k}$ and redundancy ρ_i with a single best onward lifetime metric.

C. Evaluation Metrics for Ablation Analysis

Each ablation configuration is evaluated using a detailed metric profile: PDR, deadline delivery, delay p95, energy proxy consumption, control byte overhead, and the switch activation rate. Tracking the switch rate is vital; an excessively high switch rate would indicate that the model fails to utilize the lightweight nominal branch, degrading to a heavy predictive protocol. The complete empirical results of these ablations will be documented in the final manuscript.

IX. DISCUSSION AND LIMITATIONS

A. Current Evidence and Key Insights

The empirical results of the Phase 12 evaluation strongly support the core design philosophy of GLOBE++ Lite-P+: transferring global topological intelligence into a localized routing agent via offline policy distillation. The selective Predictive Risk Switching (PRS) successfully resolves the “deployment gap” by preserving the low-overhead characteristics of geographic routing in stable environments while dynamically adapting to link breaks and routing holes when they arise. Crucially, the local student policy executes without the need for online global coordinate coordination or active multi-hop message-passing overhead.

B. Remaining Evidence Gaps and Limitations

While the current results demonstrate clear advantages, several limitations must be addressed to elevate the manuscript to submission quality:

- 1) **Full-Scale Phase 13 Validation:** Although the sub-modules of the P+ extension (link-loss gate, energy-aware tie-breaking, etc.) are implemented, their statistical verification across the full 14 scenarios and five random seeds remains a primary task.
- 2) **Physical Layer Fidelity:** The evaluations were conducted using a localized Python-based FANET simulator (Lite-GLOBE). While the simulator captures queue dynamics and link expiration rates, it simplifies multipath fading, shadow fading, and MAC-layer collisions. Cross-validation in network-level simulators like ns-3 is required to demonstrate transferability to real-world hardware.
- 3) **Absence of Traditional Ad Hoc Baselines:** The current comparison set focuses on geographic and machine learning-based baselines. Incorporating traditional reactive (AODV) and proactive (OLSR) protocols will provide a more comprehensive context for network-level overhead claims.
- 4) **Control Overhead Metrics:** The current draft reports input byte sizes as a proxy for overhead. A more rigorous approach should measure the actual control bytes transmitted per delivered packet to provide a clear trade-off comparison against message-passing algorithms like DRAMA.

C. Submission and Positioning Strategy

For high-impact SCI-level journals (e.g., IEEE Transactions on Mobile Computing or IEEE Internet of Things Journal), the paper must be positioned as a practical, deployable policy distillation method rather than a generic multi-agent RL application. The core narrative should emphasize that global graph knowledge is highly beneficial, but only when restricted to training-time privilege. By demonstrating that the resulting local decision rule is low-overhead, selective, and robust to dynamic topology shifts, the paper directly addresses the practical deployment challenges of UAV swarms.

X. CONCLUSION

This paper draft presented GLOBE++ Lite-P+, a decentralized FANET routing framework that combines global-to-local policy distillation with predictive risk switching. The fully evaluated Phase 12 predecessor, Risk-Switch Lite-GLOBE-P, achieves the strongest overall PDR and deadline delivery among the implemented comparison set while reducing input footprint relative to predictive and message-passing external baselines. The Phase 13 P+ extension introduces link-loss-aware switching, top- k onward stability, energy-aware tie breaking, and drop suppression to address the remaining weaknesses found in Phase 12.

The next required step is full Phase 13 validation across the same 14-scenario, five-seed protocol, followed by component

ablations and explicit communication-overhead reporting. Until then, Phase 12 results should be treated as verified evidence and Phase 13 P+ as the final proposed mechanism under validation.

REFERENCES

- [1] B. Karp and H. T. Kung, "Gpsr: Greedy perimeter stateless routing for wireless networks," in *Proceedings of the Annual International Conference on Mobile Computing and Networking (MobiCom)*. ACM, 2000, pp. 243–254.
- [2] W. Zhang, C. Liu, J. Jiang *et al.*, "Drama: A dynamic packet routing algorithm using multi-agent reinforcement learning with emergent communication," in *Proceedings of the International Joint Conference on Neural Networks (IJCNN)*. IEEE, 2025, pp. 1–8.
- [3] M. Xu, Y. Xia, W. Liu, and D. Huang, "Reinforcement-learning-based geographic routing considering future evolution of link states for uav networks," *IEEE Transactions on Vehicular Technology*, vol. 75, no. 2, pp. 1120–1134, 2026.
- [4] A. I. Ahmed *et al.*, "Glo-mappo: Multi-agent deep reinforcement learning for energy-efficient uav-assisted lora networks," 2025.